

# GHANENDRA SAHU

Data Scientist | ML & GenAI Engineer | Agentic AI | Python | SQL | NLP | RAG | FastAPI | Docker | AWS  
Bengaluru, India | +91 7608060384 | [ghanendra2911@gmail.com](mailto:ghanendra2911@gmail.com) | [linkedin.com/in/ghanendrasahu](https://www.linkedin.com/in/ghanendrasahu) | [github.com/ghanendrasahu](https://github.com/ghanendrasahu)

## PROFESSIONAL SUMMARY

Data Scientist with hands-on experience designing and deploying end-to-end ML inference pipelines, NLP systems, and RAG/agentic AI applications using Python, FastAPI, Docker, and AWS. Proven ability to deliver production-ready data products — from feature engineering and containerized API deployment to Power BI dashboards — with growing expertise in LLM-based agentic frameworks (LangChain, LlamaIndex, CrewAI, AutoGen), predictive analytics, real-time model serving, and operational data analysis at scale.

## TECHNICAL SKILLS

|                                    |                                                                                                                                                                                                                                    |
|------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Languages &amp; Databases:</b>  | Python, SQL (MySQL, PostgreSQL), Bash, Data Structures & Algorithms (DSA)                                                                                                                                                          |
| <b>ML / Algorithms &amp; Eval:</b> | Regression, Classification, Random Forest, XGBoost, KNN, Linear SVC, Feature Engineering, Hyperparameter Tuning, Cross-Validation, ROC-AUC, F1-Score, R-Squared, Precision-Recall, Confusion Matrix, Threshold Tuning, A/B Testing |
| <b>AI / NLP / GenAI / Agentic:</b> | NLP, TF-IDF, Semantic Search, RAG Pipelines, LLMs, Embeddings, Vector Retrieval, Text Classification, Prompt Engineering, LangChain, LlamaIndex, CrewAI, AutoGen, Multi-Agent Orchestration                                        |
| <b>Libraries:</b>                  | Pandas, NumPy, Scikit-learn, PyTorch, Matplotlib, Seaborn                                                                                                                                                                          |
| <b>Deployment &amp; MLOps:</b>     | FastAPI, Streamlit, Docker, Containerized Deployment, CI/CD, Git, GitHub, Pytest, Model Serving, Artifact Versioning                                                                                                               |
| <b>Cloud &amp; BI:</b>             | AWS (EC2, S3, Lambda), Power BI, Tableau, Excel, ETL, KPI Dashboards, Data Visualization                                                                                                                                           |

## PROFESSIONAL EXPERIENCE

**Data Science with AI Apprentice — BDreamz Global Solutions Pvt. Ltd., Bengaluru** Oct 2025 – Mar 2026  
*via Besant Technologies*

- Collected, cleaned, and organized large-scale datasets; assisted in building and validating ML/AI models under guidance of a Senior Data Scientist, collaborating on real-time data science and AI projects to identify trends and predictive insights.
- Performed data validation, preprocessing, and feature engineering; used visualization techniques to present findings with clear technical documentation.

**Data Science Intern — Princeton Smart Engineers, Bengaluru** Jan 2025 – Feb 2025

- Developed an ML pipeline for workforce attendance forecasting using Linear Regression, KNN, and Linear SVC on a 1,000-record dataset; Linear SVC achieved 89% classification accuracy — presented findings and model recommendations to senior engineers.
- Executed full preprocessing — encoding, train-test split, scaling — and evaluated models via R-Squared, F1-Score, and confusion matrix for reliable model selection.

## PROJECTS

**Customer Churn Prediction** Live Demo | GitHub

- Designed and deployed a production-grade ML inference pipeline on a 7,043-row Telco dataset (23.1% churn); engineered domain features (tenure group, charge-per-service, contract flag), achieving 87%+ accuracy with Random Forest optimized via ROC-AUC and F1-Score.
- Containerized FastAPI REST endpoint for real-time churn probability serving and a Streamlit dashboard for customer risk segmentation and automated retention recommendations.

*Tech Stack: Python, Pandas, Scikit-learn, Random Forest, FastAPI, Streamlit, Docker, REST API, Model Artifact Versioning*

**Local RAG Assistant** GitHub

- Architected an offline RAG pipeline for low-latency Q&A over enterprise TXT/PDF documents — document ingestion, chunking, TF-IDF semantic retrieval, and persistent vector indexing with zero external API dependency.
- Built a CLI query interface with top-k source-attributed results and relevance scoring for auditable document search in enterprise knowledge management workflows.

*Tech Stack: Python, RAG Architecture, TF-IDF, Semantic Search, Vector Retrieval, Persistent Indexing, PDF/TXT Parsing, CLI*

**Credit Risk & Loan Default Prediction** GitHub

- Built binary classification inference pipeline benchmarking Logistic Regression vs. Random Forest on loan features (income, credit history, employment length, interest rate); threshold tuning maximized recall on high-risk defaults, minimizing false negatives.
- Deployed a Streamlit model-serving app with real-time default probability, threshold-based decisions, and full model metadata for transparent, auditable lending operations.

*Tech Stack: Python, Pandas, Scikit-learn, Logistic Regression, Random Forest, Threshold Tuning, Streamlit, Model Serving*

## EDUCATION

**Bachelor of Computer Application (BCA) — Krupanidhi Degree College, Bangalore North University** 2022 – 2025

CGPA: 9.15/10 | Coursework: Machine Learning, Statistics, DBMS, Data Structures, Python, Data Visualization

**Data Science & ML Training (Full-Time) — Besant Technologies, Bengaluru** Sep 2025 – May 2026

## CERTIFICATIONS

- AI Apprentice — BDreamz Global Solutions / Besant Technologies
- Data Science Internship — Princeton Smart Engineers
- Python & SQL Certification — Besant Technologies
- KRUPADECON Conference Participation Certificate — Krupanidhi Group of Institutions